

Bayesian Hierarchical Modelling for Tailoring Metric Thresholds

Neil A. Ernst

Department of Computer Science – University of Victoria
nernst@uvic.ca

ABSTRACT

Software is highly contextual. While there are cross-cutting ‘global’ lessons, individual software projects exhibit many ‘local’ properties. This data heterogeneity makes drawing local conclusions from global data dangerous. A key research challenge is to construct locally accurate prediction models that are informed by global characteristics and data volumes. Previous work has tackled this problem using clustering and transfer learning approaches, which identify locally similar characteristics. This paper applies a simpler approach known as Bayesian hierarchical modeling. We show that hierarchical modeling supports cross-project comparisons, while preserving local context. To demonstrate the approach, we conduct a conceptual replication of an existing study on setting software metrics thresholds. Our emerging results show our hierarchical model reduces model prediction error compared to a global approach by up to 50%.

CCS CONCEPTS

• **Mathematics of computing** → **Bayesian computation**; • **Software and its engineering** → **Software maintenance tools**;

KEYWORDS

hierarchical models, metrics thresholds, probabilistic programming

ACM Reference Format:

Neil A. Ernst. 2018. Bayesian Hierarchical Modelling for Tailoring Metric Thresholds. In *MSR '18: MSR '18: 15th International Conference on Mining Software Repositories*, May 28–29, 2018, Gothenburg, Sweden. ACM, New York, NY, USA, 5 pages. <https://doi.org/10.1145/3196398.3196443>

1 INTRODUCTION

There are two problems with doing statistical inference on software projects. The first is that finding small *true* effects is often only possible with large datasets, yet many software projects are comparatively small. The second problem is that software characteristics (such as object oriented metrics) differ from one project to the next [25]. Posnett [21] terms this the *ecological inference problem*. As a result, many software researchers, particularly in the defect prediction and effort estimation communities, have moved to cross-project approaches, including clustering (identifying all projects with similar characteristics) and transfer learning (finding

a function that maps a classifier from one project onto the characteristics of another project). Both have underlying models that are hard to interpret, and ignore domain-specific knowledge about how software is constructed.

This paper’s new idea is to use a simpler and more understandable model for learning from software data, Bayesian hierarchical models (also known as multi-level models or random effects models [22])¹. A hierarchical model is one where the model parameters themselves have parameters drawn from a probability distribution. Parameters are things like mean and variance in a Normal distribution. Hierarchical models calculate project-specific metrics, while still taking into account the global ecological commonalities. Among other benefits, this means that inference on projects with small datasets (level 1) can be improved by our knowledge of the global set of data (level 2).

Our prediction task is to find the mean Coupling Between Objects metric for an individual project. To solve this prediction task, we compare hierarchical modeling to two other linear regression approaches. A *global* model (approach 1) pools all file information into a single level. Any interesting project-level variance is lost. A *local* or unpooled model (approach 2) calculates local-only regression coefficients. No global-level data is used, so it exhibits ‘anterograde amnesia’—any lessons from analyzing previous projects are forgotten. Projects with few data points will be untrustworthy (high standard errors). Finally, a *hierarchical*, or partial pooling model (approach 3) fits a linear model per project, but regularizing the prediction using the global data.

We find that partial pooling is much better than the other two approaches, using root mean squared error (RMSE). The advantage of the partial pooling approach:

- (1) We have an explicit prior, so it is easy to understand what is in the model and how the model is created.
- (2) Sophisticated tooling supports hierarchical models (e.g., Stan [5]), so off the shelf solutions are available.
- (3) Partial pooling accommodates our intuitive understanding of software development analysis, that balances the local and global data available.

While popular in social sciences (one author has said “multi-level modeling deserves to be the default form of regression [16, p.14]”, this form of modeling is underused in the software engineering literature. In 1998 Pickard, Kitchenham and Jones [20] pointed out that combining results in software engineering could leverage new advances in Bayesian hierarchical models. It does not seem as though there has been much take up of this call; in §5 we highlight the one or two studies that have. The main issue has been the computational power required for producing the posterior probability distributions. As we show, this is now resolved with modern inference engines and processor power.

¹‘modeling’ is used here in the statistical sense, not the software abstraction sense

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than the author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

MSR '18, May 28–29, 2018, Gothenburg, Sweden

© 2018 Copyright held by the owner/author(s). Publication rights licensed to the Association for Computing Machinery.

ACM ISBN 978-1-4503-5716-6/18/05...\$15.00

<https://doi.org/10.1145/3196398.3196443>

We investigate hierarchical models with a conceptual replication of the study of Aniche et al. [3] (henceforth SATT). The SATT study postulated that the *architectural role* a class plays (in particular, Spring framework Controllers, Views, Repositories) influences the expected value of certain software metrics. They assign a different threshold for various Chidamerer and Kemerer (CK) metric values [6] based on the role a file plays. We show how thresholds can be set more accurately using hierarchical models.

We begin by introducing some statistical background. We apply a hierarchical model to the SATT dataset [3]. We show how our model is constructed, validate it with RMSE, and discuss how such approaches might be applied in the future. An available Jupyter notebook demonstrates how this works in practice². Our emerging results show:

- (1) Hierarchical models are easy to explain and set up;
- (2) An example of using a probabilistic programming language for Bayesian data analysis;
- (3) Hierarchical models support analysis that takes into account the rich diversity of empirical data in software engineering;
- (4) Hierarchical models outperform purely local or purely global approaches.

2 BACKGROUND

For a general introduction to a Bayesian approach to statistics, people can refer to [10] and [16]. Bayesian inference is built on Bayes's theorem, $P(H|D) = \frac{P(D|H)P(H)}{P(D)}$, where H is our hypothesis (i.e., architectural role influences software metrics), and D is the data we gather. Bayesian inference calculates a *posterior probability distribution* $P(H|D)$. It does so as a consequence of the assumptions: $P(D|H)$, the *likelihood*; the parameters we wish to estimate; and $P(H)$ our *prior probability* for those parameters. We must explicitly assign a prior probability in Bayesian statistics. Bayesian inference is a machine that *conditions* on the data to generate a posterior, given the assumptions. This machine is often hard to construct mathematically, particularly in the hierarchical models we introduce in this paper. As a result, probabilistic programming techniques are needed to compute the posterior.

Probabilistic programming is a programming paradigm that uses an inference engine to fit complex and multi-level models. Variables are not deterministic but rather stochastic. Their value comes from a probability distribution. In order to compute this probability distribution, probabilistic programming languages use Markov Chain Monte Carlo (MCMC) sampling (specifically, Hamiltonian Monte Carlo and the No U-Turn Sampler). A good survey of probabilistic programming is available at [1]. One example of a probabilistic programming language (and the one used in this paper) is Stan [5], and its PyStan library (the PyMC3 package is another variant). The Stan probabilistic program to compute a linear regression $y \sim \mathcal{N}(\mu, \sigma)$ with μ , our linear model, represented by $\beta_1 + \beta_2 x_i$, looks like:

```
model { y ~ normal(beta[1] + beta[2] * x, sigma); }
data {
  int<lower=0> N;
  vector[N] x;
  vector[N] y; }
```

²<https://figshare.com/s/fd34d562ce882d1ab4d2>

```
parameters {
  vector[2] beta;
  real<lower=0> sigma; }
```

where we provide input for the data section (e.g., our vector of metric values), and Stan does the inference to find values for the parameters section that maximize the posterior, conditioned on the data.

3 METHODOLOGY

Our research question is whether a hierarchical model is more accurate than a pooled or unpooled regression model. As a working example, we perform a conceptual replication of the study of Aniche et al. [3] (SATT). This study explored how metrics thresholds could be contextualized based on architectural role. We use their data to fit three different regression models, then validate the accuracy of these models using RMSE. Once we have a model, we can use that to estimate threshold values (i.e., level at which to take action on refactoring). To identify if a file violates a threshold, one looks up its role, and then retrieves the three thresholds and compares the file's metric against the thresholds.

We focus on a narrower analysis than SATT, restricting it to the Chidamber and Kemerer (CK) metric Coupling Between Objects (CBO) [6] and the Spring Controller role. We choose CBO since the SATT results found it highly correlated to architectural role, but our basic approach is easily extended. We have 120 projects in the SATT dataset, 115 of which have Controllers, and a wide range of values for number of files. There are 56,106 files overall. Some projects are very tiny. The total file count of the smallest is 28, and the number of Controller files in that project (for example) is 18 (recall that handling such small datasets is one of the strengths of hierarchical models). We calculate the Mann-Whitney measure and Cliff's Delta measure for Controller metrics for all projects, compared to all other roles (including no role), and find a large effect (0.657), just like the original paper did.

The original study used a non-parametric test to show the effect of architectural role, but for our purposes we will need a probability distribution to serve as the likelihood. We choose a lognormal, as this is empirically accurate [13], and the same as that chosen in the SATT study [3]. We begin by taking the log of a file's CBO score (LCBO).

3.1 Modeling the Problem

We will use a simple linear regression model to estimate the posterior distribution of (L)CBO values, and do this in three different ways. We follow the model approach of Chris Fonnesbeck [8]. Fig. 2 on page 3 shows the overall comparison of the three models.

Global Pooling Global pooling aggregates all individual files and fits a single linear regression model (a fixed effects model). We predict LCBO score, y , as a function of whether the file is a controller (0/1), x : $y_i = \alpha + \beta x_i + \epsilon_i$, with equivalent statistical model expressed as $y \sim \mathcal{N}(\alpha + \beta x, \sigma)$.

What we have created is a posterior probability distribution for the unknown parameters α, β, σ . We are not getting point estimates, but rather distributions. To obtain the parameters and errors, we sample from the posterior distribution, i.e., find the mean and S.E. of the samples.

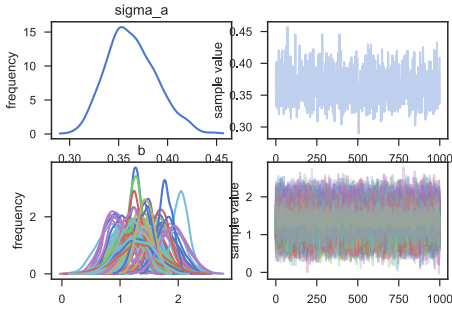


Figure 1: Checking model convergence.

Unpooled The other extreme from global pooling is to model the data per project ($N=115$). Each project gets a separate regression model, and we can compare the co-efficients for each one. The regression model is $y_i = \alpha_{i[j]} + \beta x_i + \epsilon_i$ for projects j .

Here we are favoring individual project variance. Small samples will greatly bias our project-level estimates. McIlreath [16] calls this “anterograde amnesia” since our model machinery ‘forgets’ what it learned from the previous projects.

Partial Pooling - Varying Slope and Intercept The hierarchical, *partial pooling* approach adjusts the local estimate *conditional* on the global estimate. The hierarchy in our hierarchical model is files within projects. Partial pooling gives nuance we know exists because of local effects from architectural role [3], but *regularizes* that by the global parameters (another name for this is *shrinkage*).

The partial pooling model will allow the regression model’s slope and intercept to vary. If our pooled regression equation predicts y_i , an observation at point i , to be $y_i = \alpha + \beta x_i + \epsilon_i$, then the partial pooling regression equation is $y_i = \alpha_{j[i]} + \beta_{j[i]} x_i + \epsilon_i$. The main difference with unpooled models is that we assign hyperpriors to our parameters.

The model component of the probabilistic program is

```
model {
  mu_a ~ normal(0, 100); # hyperprior
  mu_b ~ normal(0, 100); # hyperprior
  a ~ normal(mu_a, sigma_a); # prior
  b ~ normal(mu_b, sigma_b); # prior
  # likelihood:
  y ~ normal(a[project] + b[project]*x, sigma);
}
```

with parameters a, b , priors for y , having posteriors that are normal, with priors for the mean distributed as another normal model with mean 0 and sigma 100. This is an uninformative prior (allowing for a wide range of values, and thus dominated by the data). Choosing priors effectively is a core skill in probabilistic programming.

4 RESULTS

An important part of Bayesian hierarchical modeling is model checking. This involves inspection of the model inferences to ensure the fit is reasonable. To do this we plot (Fig. 1) the posterior distribution of our parameters and the individual samples.




Figure 2: Comparing three models. Blue=unpooled, dotted black=partial pool, dashed red=pooled. N indicates number of controller files. X-axis reflects predictor of controller/not-controller.

The left side shows our marginal posterior—for each parameter value on the x-axis we get a probability on the y-axis that tells us how likely that parameter value is. The right side shows the sampling chains. Our sampling chains for the individual parameters (left side) seem well converged and stationary (there are no large drifts or other odd patterns) [22].

A sample of four projects in the SATT dataset is shown in Fig. 2. Conceptually, the regression is more accurate if the right side data points (in blue dots) are closer to the regression line ($x=1$ being files which are Controllers). The dashed black line is the partial pooling estimate; the dashed red line is the global/pooled estimate; and the solid blue line the unpooled estimate. The global estimate seems way off. Depending on the number of data values, the partial and unpooled estimates are fairly close. The improvement in accuracy of the partial pooling approach increases as the number of datapoints decreases (and the unpooled model gets worse).

Validation—We validate accuracy using a simple root mean squared error (RMSE) approach. We fit our three regression models, then calculate the average distance from the predicted score to the actual file CBO values. We expect the partial-pooling to have a lower RMSE value if it is more accurate than the others.

Table 1 reports the mean RMSE over all projects with Controllers ($n=115$), and approximate Stan model training time with 2 chains of 1000 iterations (model compilation (to C++) is a constant 30s or so in addition). Partial pooling is **nearly twice as effective** as full pooling, but takes 4 times as long (an OLS approach is nearly instantaneous and produces the same RMSE as the full pooling model). No pooling is nearly as effective as partial pooling (likely because the number of files that are *not* controllers is fairly large). However, error values for the unpooled approach vary; in projects with small numbers of files, the RMSE for an unpooled approach is much higher. This represents the benefits of shrinking towards the global mean. For example, in the Netvogue project of Fig. 2, the unpooled regression (blue) is quite far from the partial pooling line (black dashed). This is a difference in RMSE of 1.23 vs. 0.91.

Model	RMSE	Sampling Time
Full Pool	1.099	26.9s
Unpooled	0.630	5m11s
Partial Pool	0.523	3m18s

Table 1: RMSE for 3 Regression Approaches

Justifying Thresholds—The SATT finding [3] was that a file’s architectural role played a significant part in its eventual CBO numbers. They used the procedure of Alves et al. [2] to identify thresholds. The Alves thresholds were set by ranking files (across all projects) by lines of code (LOC), then finding the metric value of the file closest to 70/80/90% (moderate/high/very high risk).

We use our partial pooling regression model to set thresholds at 70/80/90% of the normal distribution of log CBO. E.g., for the project ‘V2V’, the partial pooling model tells us the project’s mean LCBO value $y \sim N(a_{\text{project}} + b_{\text{project}} * x, \sigma)$. This normal posterior distribution is used to find the point probability for our thresholds τ (i.e. the point on the normal distribution holding $\tau\%$ of the probability mass). $\exp(\tau)$ produces CBO score. For example, the ‘V2V’ project has expected Controller CBO thresholds of 32,49, and 88. Compare this to another project, ‘tatami-team’, with 21,32,58, respectively. Our thresholds vary based on the project characteristics and the regularizing global dataset. Full results for all projects are in our replication package (footnote p. 2).

5 DISCUSSION AND RELATED WORK

The chief threats to validity in Bayesian hierarchical modeling come from poor understanding of the underlying probabilistic program, including what priors to choose. Model checking is vital to mitigate this. There are subtleties in how the sampler works that need careful checking, as shown by Wiecki’s blog post [24].

Other threats to validity include data analysis and implementation issues (internal validity). One threat that is mitigated by this approach is external validity, since a hierarchical model inherently accounts for inter-project differences. Many of the common complaints about lack of generalizability (e.g., of Microsoft results to other companies) could be avoided with a hierarchical approach.

This paper has only hinted at the uses for hierarchical modeling. Other useful aspects include the ability to handle correlations between local and global variables [16], innately accounting for multiple-comparisons problems [11], and adding group-level predictors (for example, using number of files per project). Other aspects to explore include adding additional predictors such as lines of code, different underlying prior distributions, and improved analysis of model fit using Leave One Out (LOO) or Widely Applicable Information Criterion (WAIC) [23]. Although consumer-ready, Bayesian inference remains unexplored compared to the highly usable frequentist packages. Probabilistic programming is also an area that merits more investigation as a form of software development in its own right.

Related Work—There are two main categories of *related work*. First, there is work that identifies the **differences between per-project and pooled predictors** in analysing software metrics. Menzies et al. [17] introduced the concept of ‘local’ (clusters of data

with similar characteristics) and ‘global’ (cross-project, pooled) lessons in defect prediction. This idea has been followed up by many others (e.g., [15, 4, 19]). Posnett et al. [21] conducted a study comparing defect prediction at two levels: aggregated (package) vs. unaggregated (file) and introduced the concept of ecological inference. The innovation in our paper is to propose a partial approach, regularizing the local predictions with the global information. Our paper does not, however, make any claim to improve on defect prediction—we leave that to future work. Several papers [4, 26] have leveraged linear (continuous response) and logistic regression (categorical response); however, these are single-level models. Bettenburg et al. [4] blend the regression levels using a different approach, Multivariate Adaptive Regression Splines (MARS). The MARS approach appears to be fairly uncommon in the statistics literature, and uses hinge functions to reduce the residual error. The Stan models we introduce are simpler to specify, and the Bayesian approach to conditioning on the prior data we find more intuitive.

JÄyrgenson et al. [14] introduced a technique called ‘regression to the mean’ (RTM). It also attempts to regularize local predictions using global parameters. Hierarchical models are a more generalized and robust version of this, at the expense of computation time.

There are few software engineering studies using hierarchical models. The two we have identified are the 2012 work of Ehrlich and Cataldo [7], who used multi-level (hierarchical) models to assess communication in global software development. In 2017 Hassan et al. [12] used a multi-level model to study app store reviews. This approach is similar to this paper but uses a frequentist, maximum likelihood approach, i.e., not Bayesian, using the ‘lme4’ R package. We believe the Bayesian approach better accounts for prior information that is helpful when lacking observations.

Secondly, there is work that looks at **metric thresholds**. Since metrics were proposed for software, researchers have sought to identify what thresholds should trigger warnings. The paper we replicate from Aniche et al. [3] has a good survey. We point out the work of Herraiz et al. [13], Oliveira et al. [18], Foucault et al. [9] and Alves et al. [2] as more recent examples. The metrics industry, e.g., SIG and CAST, has also long used databases of company performance (a global pool) as benchmarks for evaluating new clients.

6 CONCLUSION

In this paper we introduced a hierarchical approach to setting metric thresholds. Using a partial pooling approach, we are able to account for global context, while retaining local variance. Our linear regression model defined a project-specific mean LCBO score, which we used to set the benchmarks accordingly. The hierarchical model is simple to describe, handles cross-project comparisons, and is more accurate than a purely global or purely local approach.

7 ACKNOWLEDGEMENTS

We are grateful to the authors of [3], who made their well-packaged and documented data and code available for others.

REFERENCES

- [1] Mohammed AlQuraishi. 2015. The state of probabilistic programming. (2015). <https://moalquraishi.wordpress.com/2015/03/29/the-state-of-probabilistic-programming/>.
- [2] Tiago L Alves, Christiaan Ypma, and Joost Visser. 2010. Deriving metric thresholds from benchmark data. In *Proceedings of the IEEE International Conference on Software Maintenance*, 1–10.
- [3] Mauricio Aniche, Christoph Treude, Andy Zaidman, Arie van Deursen, and Marco Aurelio Gerosa. 2016. SATT: tailoring code metric thresholds for different software architectures. In *Proceedings of the IEEE International Working Conference on Source Code Analysis and Manipulation*. (Oct. 2016).
- [4] N. Bettenburg, M. Nagappan, and A. E. Hassan. 2012. Think locally, act globally: improving defect and effort prediction models. In *IEEE Working Conference on Mining Software Repositories*. (June 2012).
- [5] Bob Carpenter et al. 2017. Stan: a probabilistic programming language. *Journal of Statistical Software*, 76, 1.
- [6] Shyam R. Chidamber and Chris F. Kemerer. 1991. Towards a metrics suite for object oriented design. *Proceedings of the ACM SIGPLAN Conference on Object-Oriented Programming, Systems, Languages, and Applications*.
- [7] Kate Ehrlich and Marcelo Cataldo. 2012. All-for-one and one-for-all?: a multi-level analysis of communication patterns and individual performance in geographically distributed software development. In *Proceedings of the ACM Conference on Computer Supported Cooperative Work*. Seattle, Washington, USA, 945–954.
- [8] Chris Fonnesbeck. 2016. A primer on bayesian multilevel modeling using pystan. (2016). https://github.com/fonnesbeck/stan_workshop_2016/blob/master/notebooks/Multilevel%20Modeling.ipynb.
- [9] Matthieu Foucault, Marc Palyart, Jean-Rémy Falleri, and Xavier Blanc. 2014. Computing contextual metric thresholds. In *Proceedings of the ACM SIGAPP Symposium on Applied Computing*.
- [10] Andrew Gelman, John B. Carlin, Hal S. Stern, David B. Dunson, Aki Vehtari, and Donald B. Rubin. 2014. *Bayesian Data Analysis*. (3rd ed.). CRC Press.
- [11] Andrew Gelman, Jennifer Hill, and Masanao Yajima. 2012. Why we (usually) don't have to worry about multiple comparisons. *Journal of Research on Educational Effectiveness*, 5, 189–211.
- [12] Safwat Hassan, Chakkrit Tantithamthavorn, Cor-Paul Bezeemer, and Ahmed E. Hassan. 2017. Studying the dialogue between users and developers of free apps in the google play store. *Proceedings of the International Symposium on Empirical Software Engineering and Measurement*, (Sept. 2017).
- [13] Israel Herraiz, Daniel Rodriguez, and Rachel Harrison. 2012. On the statistical distribution of object-oriented system properties. In *International Workshop on Emerging Trends in Software Metrics (WETSoM)*. IEEE, (June 2012).
- [14] Magne Jørgensen, Ulf Indahl, and Dag Sjøberg. 2003. Software effort estimation by analogy and “regression toward the mean”. *Journal of Systems and Software*, 68, 3, 253–262.
- [15] Yasutaka Kamei, Takafumi Fukushima, Shane McIntosh, Kazuhiro Yamashita, Naoyasu Ubayashi, and Ahmed E. Hassan. 2016. Studying just-in-time defect prediction using cross-project models. *Empirical Software Engineering*, 21, 5, (Oct. 2016), 2072–2106. ISSN: 1573-7616.
- [16] Richard McIlreath. 2016. *Statistical Rethinking: A Bayesian Course with Examples in R and Stan*. CRC Press.
- [17] T. Menzies, A. Butcher, A. Marcus, T. Zimmermann, and D. Cok. 2011. Local vs. global models for effort estimation and defect prediction. In *Proceedings of the IEEE/ACM International Conference on Automated Software Engineering*, 343–351.
- [18] P. Oliveira, M. T. Valente, and F. P. Lima. 2014. Extracting relative thresholds for source code metrics. In *Proceedings of the European Conference on Software Maintenance and Reengineering*, 254–263.
- [19] A. Panichella, R. Oliveto, and A. De Lucia. 2014. Cross-project defect prediction models: l'union fait la force. In *Proceedings of the European Conference on Software Maintenance and Reengineering*, 164–173.
- [20] Lesley M. Pickard, Barbara A. Kitchenham, and Peter W. Jones. 1998. Combining empirical results in software engineering. *Information and Software Technology*, 40, 14, 811–821.
- [21] Daryl Posnett, Vladimir Filkov, and Premkumar Devanbu. 2011. Ecological inference in empirical software engineering. In *Proceedings of the IEEE/ACM International Conference on Automated Software Engineering*, 362–371.
- [22] Tanner Sorensen, Sven Hohenstein, and Shravan Vasishth. 2016. Bayesian linear mixed models using stan: a tutorial for psychologists, linguists, and cognitive scientists. *The Quantitative Methods for Psychology*, 12, 3, 175–200.
- [23] Aki Vehtari, Andrew Gelman, and Jonah Gabry. 2017. Practical bayesian model evaluation using leave-one-out cross-validation and WAIC. *Statistics and Computing*, 27, 1413–1432, 5.
- [24] Thomas Wiecki. 2017. Why hierarchical models are awesome, tricky, and Bayesian. (2017). <http://twiecki.github.io/blog/2017/02/08/bayesian-hierarchical-non-centered/>.
- [25] F. Zhang, A. Mockus, Y. Zou, F. Khomh, and A. E. Hassan. 2013. How does context affect the distribution of software maintainability metrics? In *Proceedings of the IEEE International Conference on Software Maintenance*, 350–359.
- [26] M. Zhou and A. Mockus. 2015. Who will stay in the FLOSS community? modeling participant's initial behavior. *IEEE Transactions on Software Engineering*, 41, 1, 82–99.